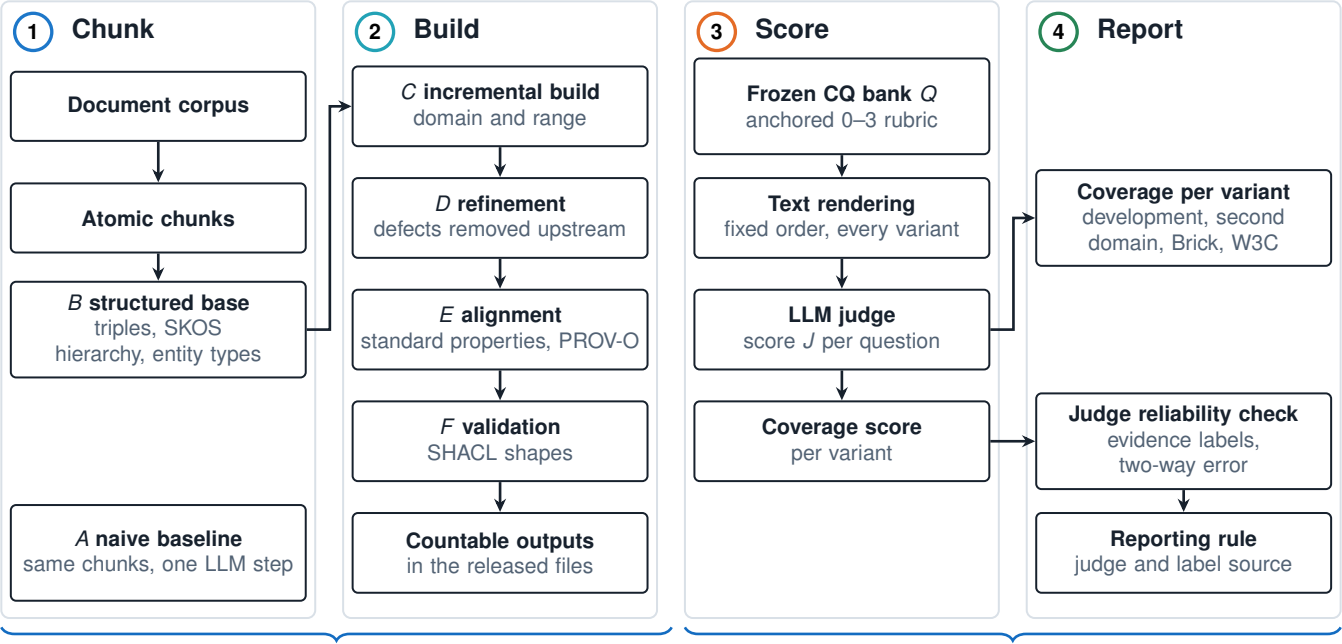
1 Chunk
Document corpus
Atomic chunks
B structured base
triples, SKOS
hierarchy, entity types
A naive baseline
same chunks, one LLM step
2 Build
C incremental build
domain and range
D refinement
defects removed upstream
E alignment
standard properties, PROV-O
F validation
SHACL shapes
Countable outputs
in the released files
3 Score
Frozen CQ bank Q
anchored 0–3 rubric
Text rendering
fixed order, every variant
LLM judge
score J per question
Coverage score
per variant
4 Report
Coverage per variant
development, second
domain, Brick, W3C
Judge reliability check
evidence labels,
two-way error
Reporting rule
judge and label source
Pipeline: every stage adds one standards function that can be counted without a judge
Evaluation: bank and rubric frozen before any artifact is scored; each score carries its judge